



---

## Lab 03: Implementing Data Pipelines with Fabric Data Factory

---

### Introduction

In this lab, you will ingest parquet containing sales data directly to the Managed Lakehouse table, using a Data Pipeline with Copy activity, to append this dataset to the existing Sales table.

### Objectives

After completing this lab, you will be better able to:

1. Create a data pipeline
2. Execute and monitor the pipeline
3. Create a basic Semantic Model
4. Auto-create a Power BI report to visualize data

### Estimated time to complete this lab

60 minutes

Contents

Lab 03: Implementing Data Pipelines with Fabric Data Factory..... 1

    Introduction ..... 1

    Objectives ..... 1

    Task 1: Create a Data pipeline ..... 3

    Task 2: Execute and monitor the pipeline..... 8

    Task 3: Create the Semantic Model..... 11

    Task 4: Create the Report ..... 15

Lab Prerequisites

- Workspace: Fabric, Power Premium or Fabric trial
- Individual license: Power Pro or Premium Per User account

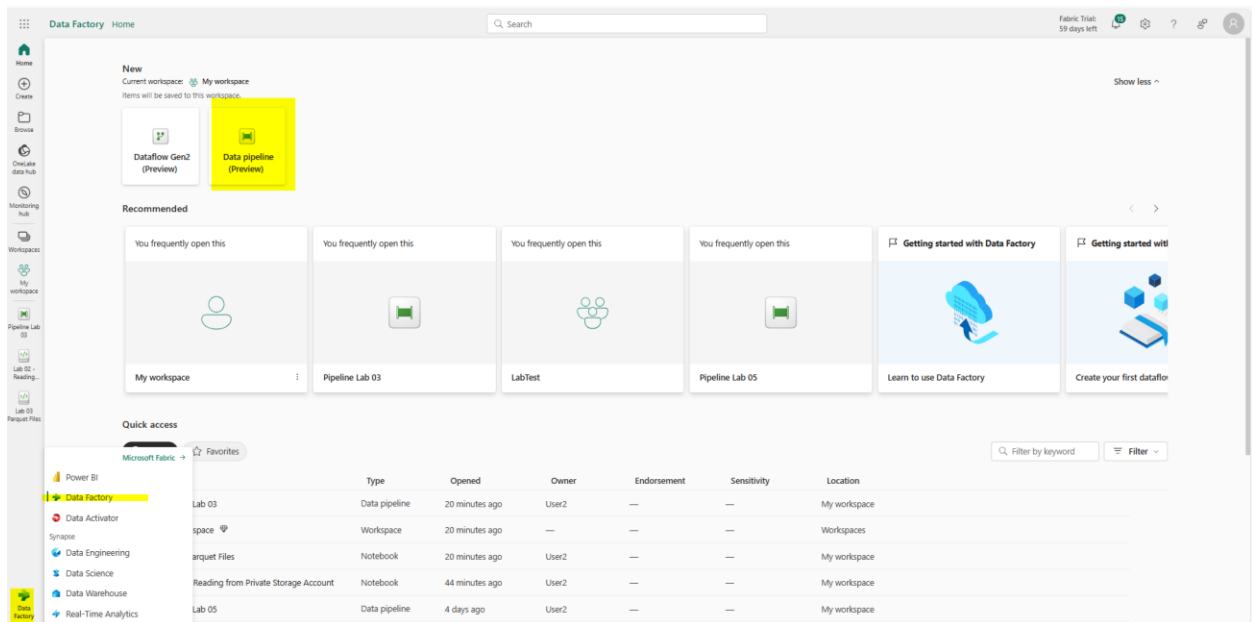
Information provided by your training provider

- Trial tenant (if applicable): login & password, workspace to use for the lab.
- Azure Data Lake Gen2 (containing data sources): account name & shared access signature.


## Task 1: Create a Data pipeline

In this task, you will create the Copy Activity that copies the parquet files from a private storage account to the Sales table from the Managed Lakehouse.

- On the MS Fabric menu, navigate to Data Factory and click on the Data pipeline. Provide the name Pipeline Lab 03

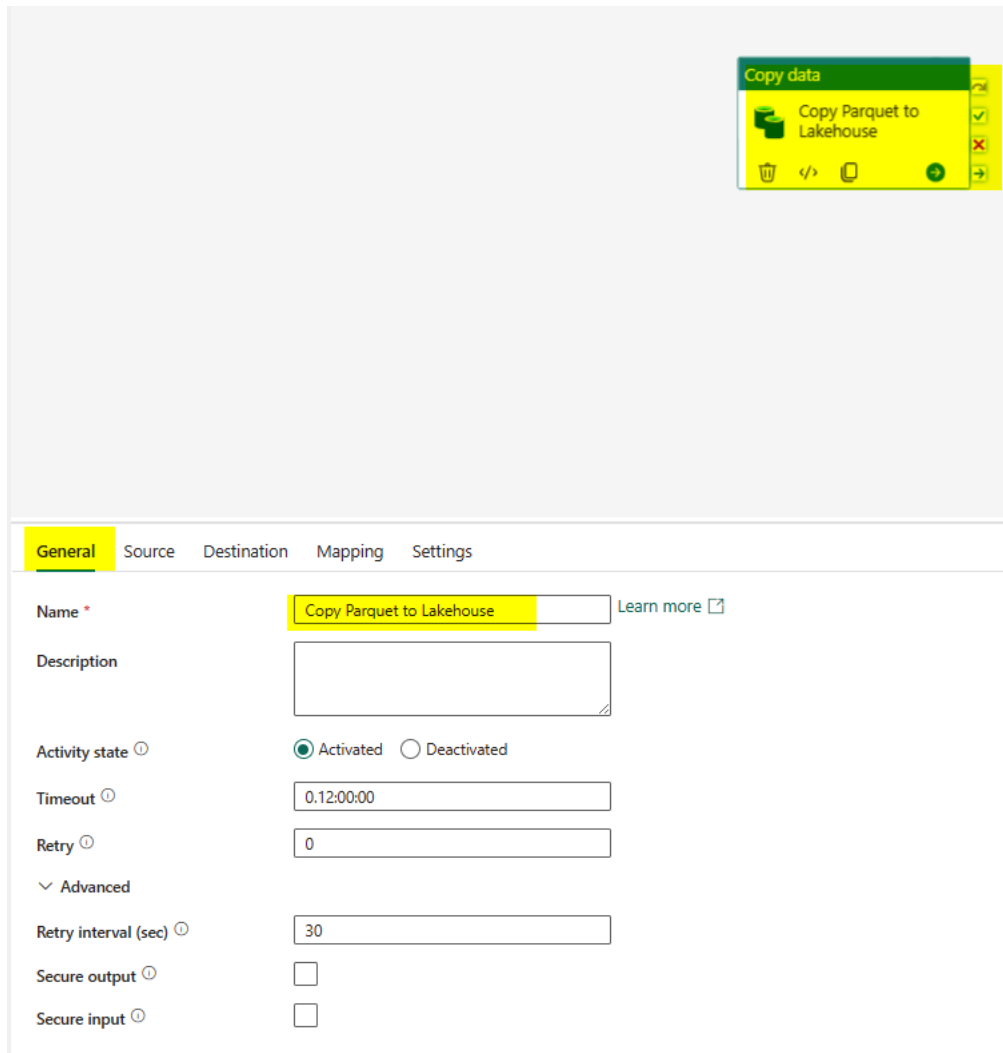


- From the top menu, click on the Copy data and select Add to canvas:



- Click on the Copy data that you just added to the Canvas:

Go to the **General** Tab and under the Name type: Copy Parquet to Lakehouse. The other settings under this tab will remain unchanged.



The screenshot shows the 'Copy data' activity configuration interface in Microsoft Fabric Data Factory. The 'General' tab is active, displaying the following settings:

- Name:** Copy Parquet to Lakehouse (with a 'Learn more' link)
- Description:** (empty text box)
- Activity state:** ☒ Activated, ☐ Deactivated
- Timeout:** 0.12:00:00
- Retry:** 0
- Advanced:**
  - Retry interval (sec):** 30
  - Secure output:** ☐
  - Secure input:** ☐

A floating 'Copy data' dialog is visible in the top right corner, showing the activity name 'Copy Parquet to Lakehouse' and various action icons.

- Click on Source Tab and select the following:
  - Data store type: External
  - Connection: Choose the name you provided in the first lab when you connected to the Private Storage account and test the connection to be successful
  - File path type: File path
  - File path: click on browse and navigate to the container labdata/parquet/sales
  - **Recursively:** checked ( this will ensure that all the files from the subfolders labdata/parquet/sales will be read)
  - File format: Parquet
  - All the other settings from this tab will remain unchanged.

Make sure that your Source tab looks like this at the end:

## Upskilling on MS Fabric Data Factory

The screenshot shows the 'Source' tab of the MS Fabric Data Factory configuration. The 'Data store type' is set to 'External'. The 'Connection' is 'fabricdatafactorylab\_adlsgen2'. The 'File path type' is 'File path'. The 'File path' is 'labdata / parquet/sales'. The 'File format' is 'Parquet'. The 'Advanced' section is expanded, showing 'Filter by last modified', 'Enable partitions discovery', 'Max concurrent connections', and 'Additional columns'.

- Click on the Destination tab:
  - Data store type: Workspace
  - Workspace data store type: Lakehouse
  - Lakehouse: Managed( this should be the name you provided on the first lab)
  - Root folder: Tables
  - Table name: Sales
  - Expand Advanced
    - – Table action: Append

At the end the Destination tab should look like this:

The screenshot shows the 'Destination' tab of the MS Fabric Data Factory configuration. The 'Data store type' is 'Workspace'. The 'Workspace data store type' is 'Lakehouse'. The 'Lakehouse' is 'Managed'. The 'Root folder' is 'Tables'. The 'Table name' is 'Sales'. The 'Advanced' section is expanded, showing 'Table action' set to 'Append'.

- On the Mapping tab, the mapping should be done automatically. If not, click on the Import schemas. At the end the Mapping tab should look like this:

General Source Destination **Mapping** Settings

> Type conversion settings

Import schemas Preview source

+ New mapping Reset Delete

<input type="checkbox"/>	Source	Type	→	Destination	Type		
<input type="checkbox"/>	SalesOrderNumber	abc STRING	→	SalesOrderNumber	string	+	🗑️
<input type="checkbox"/>	SalesOrderLineNumber	abc STRING	→	SalesOrderLineNumber	integer	+	🗑️
<input type="checkbox"/>	OnlineSalesKey	abc STRING	→	OnlineSalesKey	integer	+	🗑️
<input type="checkbox"/>	OrderDate	abc STRING	→	OrderDate	timestamp	+	🗑️
<input type="checkbox"/>	DeliveryDate	abc STRING	→	DeliveryDate	timestamp	+	🗑️
<input type="checkbox"/>	StoreKey	abc STRING	→	StoreKey	integer	+	🗑️
<input type="checkbox"/>	ProductKey	abc STRING	→	ProductKey	integer	+	🗑️
<input type="checkbox"/>	PromotionKey	abc STRING	→	PromotionKey	integer	+	🗑️
<input type="checkbox"/>	CustomerKey	abc STRING	→	CustomerKey	integer	+	🗑️
<input type="checkbox"/>	SalesQuantity	abc STRING	→	SalesQuantity	integer	+	🗑️
<input type="checkbox"/>	SalesAmount	abc STRING	→	SalesAmount	string	+	🗑️
<input type="checkbox"/>	ReturnQuantity	abc STRING	→	ReturnQuantity	integer	+	🗑️
<input type="checkbox"/>	ReturnAmount	abc STRING	→	ReturnAmount	string	+	🗑️
<input type="checkbox"/>	DiscountQuantity	abc STRING	→	DiscountQuantity	integer	+	🗑️
<input type="checkbox"/>	DiscountAmount	abc STRING	→	DiscountAmount	string	+	🗑️
<input type="checkbox"/>	TotalCost	abc STRING	→	TotalCost	string	+	🗑️
<input type="checkbox"/>	UnitCost	abc STRING	→	UnitCost	string	+	🗑️
<input type="checkbox"/>	UnitPrice	abc STRING	→	UnitPrice	string	+	🗑️

- The Settings tab remains unchanged:

General Source Destination Mapping **Settings**

Intelligent throughput optimization ⓘ Auto

☐ Use custom value

Degree of copy parallelism ⓘ Auto

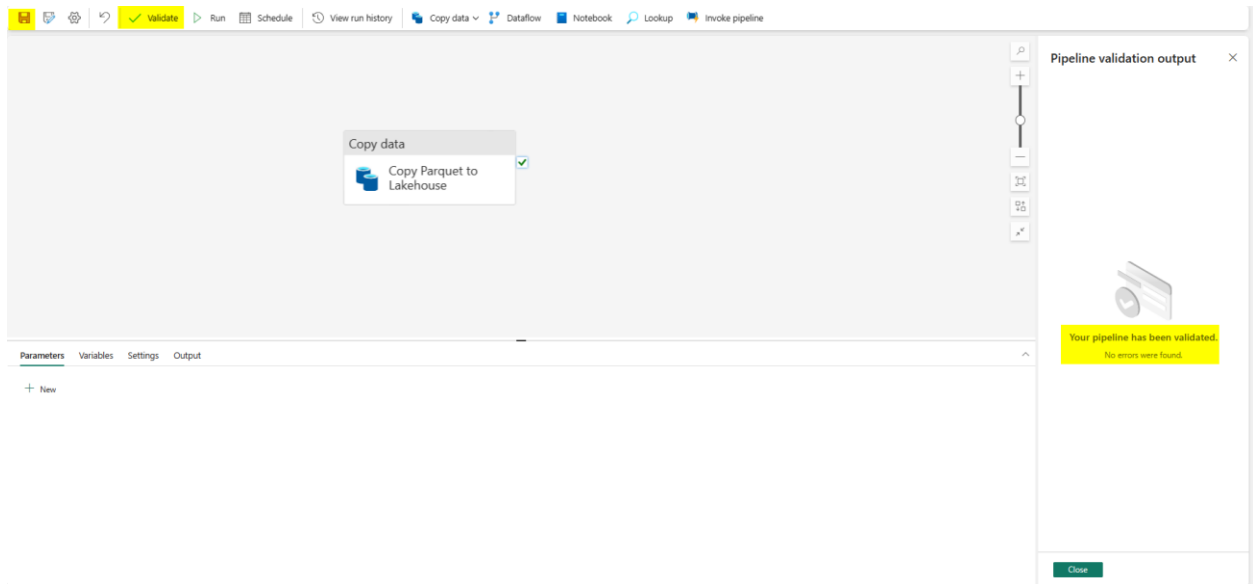
Fault tolerance ⓘ

Enable logging ⓘ ☐

Enable staging ⓘ ☐

- Once you finish building the Pipeline, from the top tab, click on the Validate to make sure that there are no errors and then finally on the Save button.

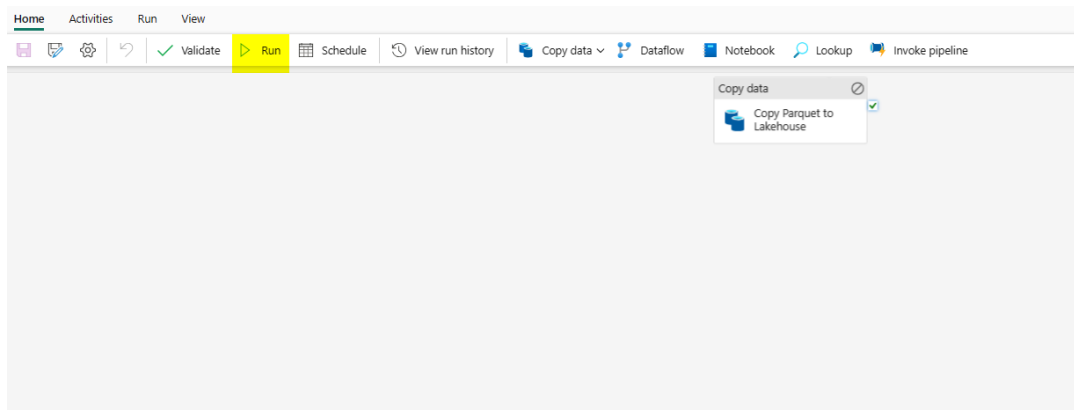
## Upskilling on MS Fabric Data Factory



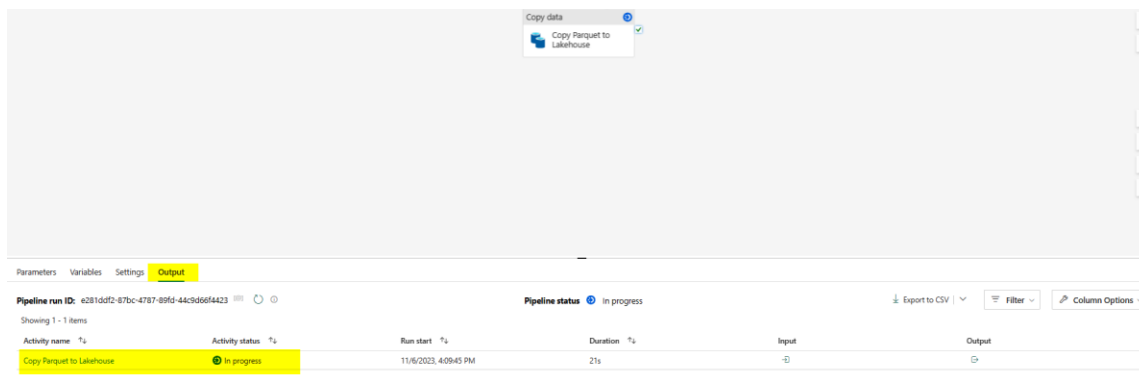
## Task 2: Execute and monitor the pipeline

In this task, you will start by trigger the pipeline

- From the top menu click on the Run button.

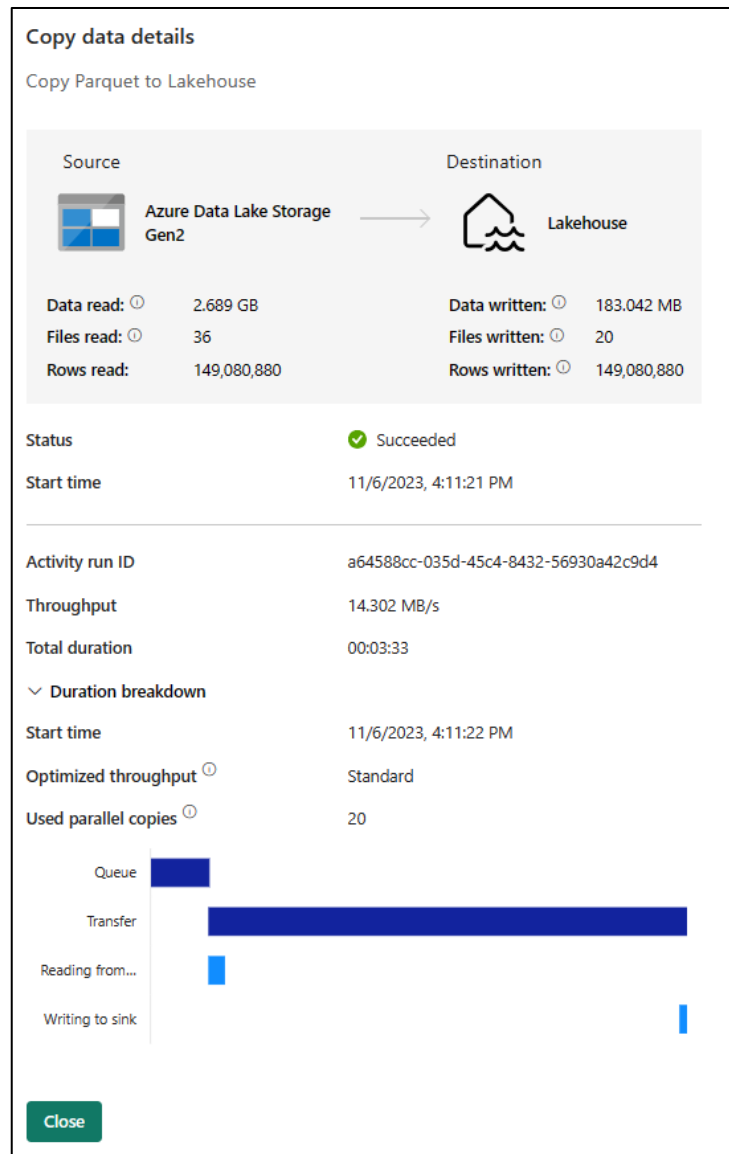


Then at the bottom of the page click on the Output and then on the Pipeline:



You can explore the copy data details to understand how much data you have transferred to the Lakehouse table.

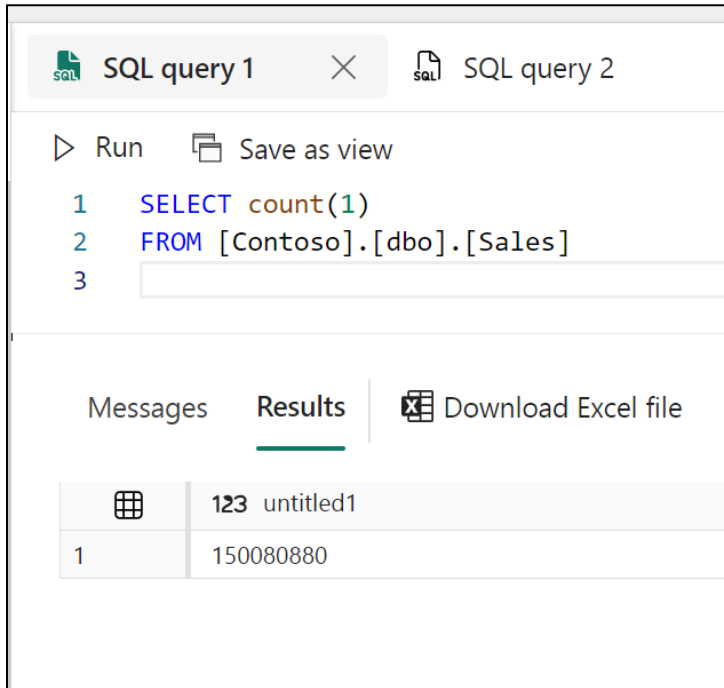




As it can be observed, at the time we read the data from ADLS Gen 2 we had 2.689GB with 36 files, however, upon migration to the Lakehouse, the dataset transformed into a more compact form, occupying 183.042MB distributed across 20 files

This exercise illustrates the efficiency of transforming data from ADLS Gen 2 into a Delta table within the Lakehouse environment. By leveraging Delta tables, data compression is improved, and smaller files are consolidated into larger ones, optimizing data storage and access.

- Come back to the SQL Analytics Endpoint of your Lakehouse, and run this basic SQL query to count the number of rows.
- You can realize that the Sales now contains more than 150M of rows.



The screenshot shows the SQL query editor interface. At the top, there are tabs for 'SQL query 1' and 'SQL query 2'. Below the tabs, there are buttons for 'Run' and 'Save as view'. The query text is as follows:

```
1 SELECT count(1)
2 FROM [Contoso].[dbo].[Sales]
3
```

Below the query editor, there are tabs for 'Messages', 'Results', and 'Download Excel file'. The 'Results' tab is selected, showing a table with 123 rows. The first row contains the value 150080880.

	123 untitled1
1	150080880

- You can also use the previously created SQL View to control sales per store.



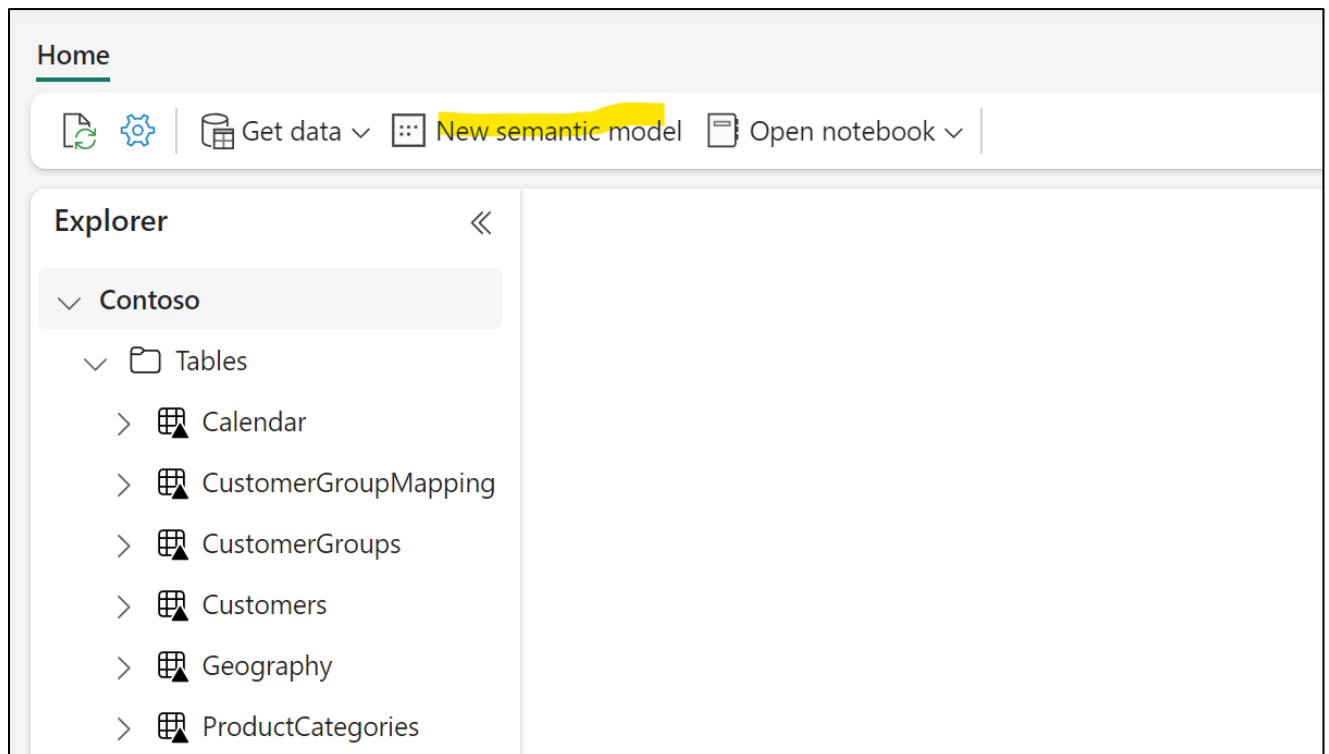
The screenshot shows the 'Data preview' table in the SQL Analytics Endpoint. The table has three columns: 'StoreName', 'NbRows', and 'SalesQuantity'. The table is showing 3 rows of data.

	ABC StoreName	123 NbRows	123 SalesQuantity
1	Contoso Asia Online Store	46826387	46437912
2	Contoso North America Online Store	56548536	56103800
3	Contoso Europe Online Store	46705957	46281994

## Task 3: Create the Semantic Model

In this task, you will create a simple Semantic Model from the Lakehouse, based on the Direct Lakehouse Mode.

- From the Lab Workspace, select the Contoso Lakehouse
- In the Lakehouse viewer, click on the New semantic model button



- On the New semantic model UI, type SalesAnalysis as semantic model name, and select the following tables :
  - Sales
  - Calendar
  - Customers
  - Geography
  - Products
  - Stores
  - Promotions
- Click on Submit to initiate the Dataset creation

### New semantic model

Name

SalesAnalysis

Workspace

LabAdmin

Select or deselect objects for the semantic model. Only objects that can be added to the semantic model are shown. [Learn more](#)

Q Search

☐ Select all

☒ Sales

☒ Calendar

☐ CustomerGroupMapping

☐ CustomerGroups

☒ Customers

☒ Geography

☐ ProductCategories

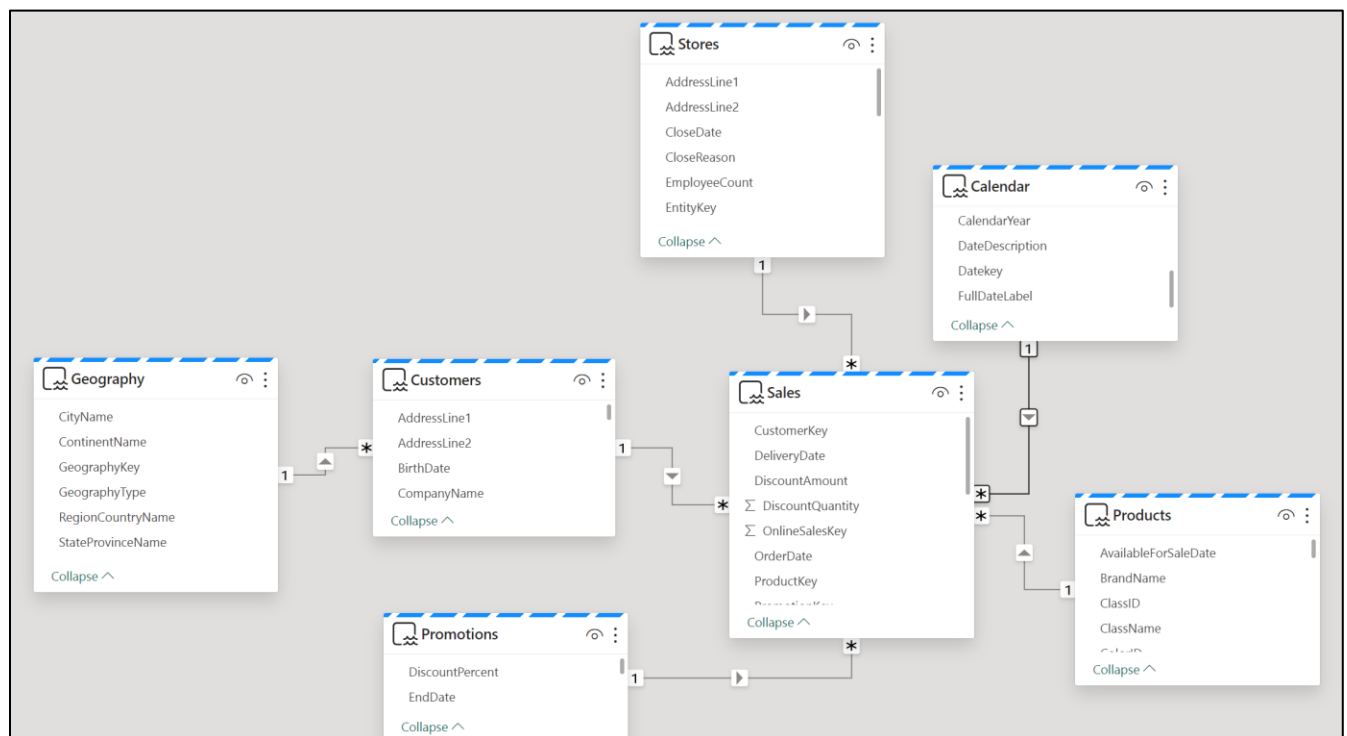
☐ ProductSubCategories

☒ Products

Confirm

Cancel

- The objective is now to design the semantic model with the following relationships :



- Click on the Manage relationships button to create validate the required relationships between tables.

Manage relationships

+ New relationship

Edit

Delete

Filter

<input type="checkbox"/>	From: table (column) ↑	Relationship	To: table (column)	Status	
<input type="checkbox"/>	Customers (GeographyKey)		Geography (GeographyKey)	Active	...
<input type="checkbox"/>	Sales (CustomerKey)		Customers (CustomerKey)	Active	...
<input type="checkbox"/>	Sales (OrderDate)		Calendar (Datekey)	Active	...
<input type="checkbox"/>	Sales (ProductKey)		Products (ProductKey)	Active	...
<input type="checkbox"/>	Sales (PromotionKey)		Promotions (PromotionKey)	Active	...
<input type="checkbox"/>	Sales (StoreKey)		Stores (StoreKey)	Active	...

- You will now implement 2 very simple Measures (business calculations)
  - Create the measures Quantity using the following DAX definition:
- Quantity=SUM(Sales[SalesQuantity])

Home Help

New measure

New column

New table

Calculation group

Manage roles

Manage relationships

New report

Edit tables

Calculations

Security

Relationships

Reporting

Modeling

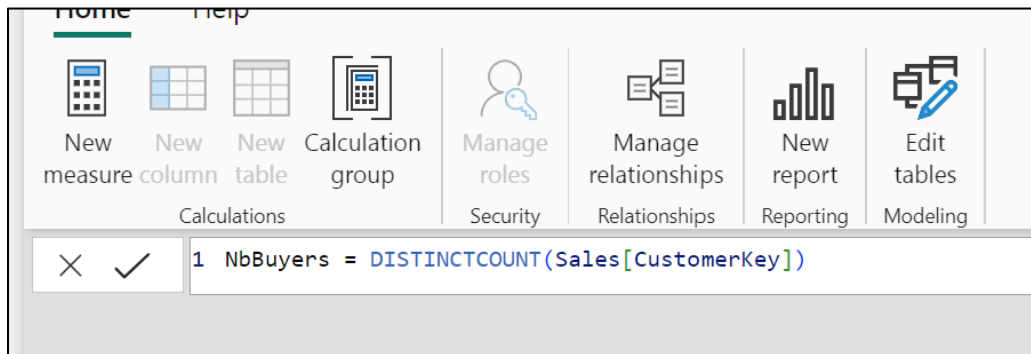
×

✓

1 Quantity = SUM(Sales[SalesQuantity])

- Create the measures NbBuyers using the following DAX definition:

NbBuyers=**DISTINCTCOUNT**(Sales[CustomerKey])



The semantic model is now ready to support report visualizations.

## Task 4: Create the Report

In this task, you will create a simple report to visualize data from the dataset using the Auto-create report function, then you will customize it.

- From the Lab workspace, select the SalesAnalysis semantic mode, click on the dots and select the option Auto-create report

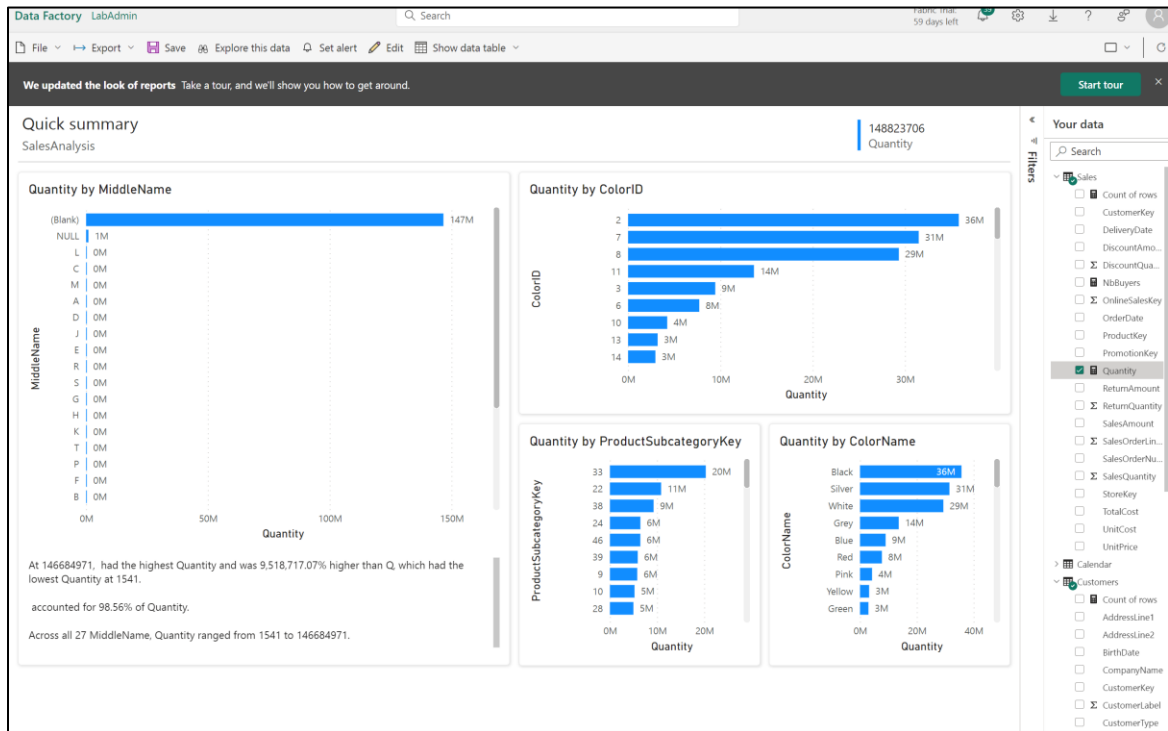
The screenshot shows the LabAdmin interface with a list of resources. The 'SalesAnalysis' semantic model is selected, and a context menu is open with 'Auto-create report' highlighted.

Name	Type	Owner
Contoso	Lakehouse	VBD Admin
Contoso	Semantic model (...)	LabAdmin
Contoso	SQL analytics end...	LabAdmin
Contoso	Data pipeline	VBD Admin
Lab 02A - Load Tables from CSV	Notebook	VBD Admin
Lab 02B - Reading from Private Storage Account	Notebook	VBD Admin
Pipeline Lab 03	Data pipeline	VBD Admin
SalesAnalysis	Semantic model	LabAdmin

Context menu options for 'SalesAnalysis':

- Explore this data (preview)
- Analyze in Excel
- Create report
- Auto-create report**
- Create paginated report
- Delete

- After few seconds, a basic report is created



- Then you can personalize each visual by changing the Visualization type and the columns and measures displayed on the visual axis.

**Personalize**

147M

Visualization type  
Clustered bar chart

Y-axis  
MiddleName

X-axis  
Quantity

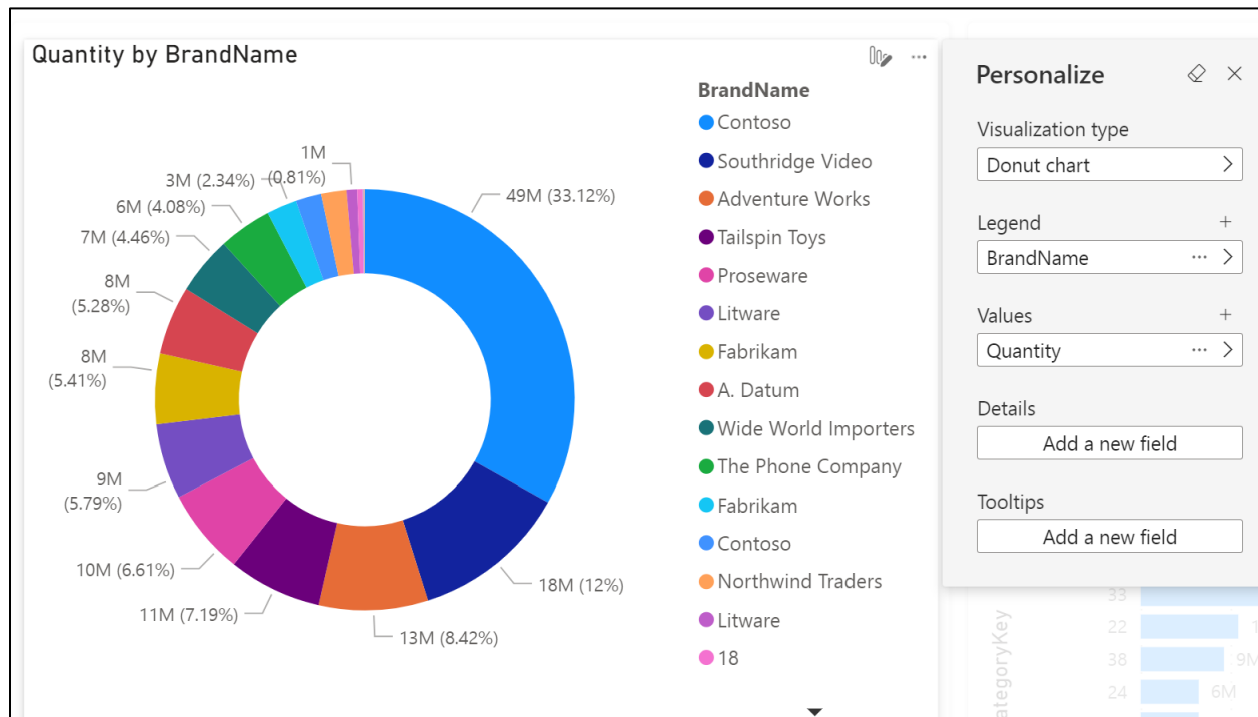
Legend  
Add a new field

Small multiples  
Add a new field

Tooltips  
Add a new field



- The screenshot hereafter displays the result using a donut chart with Sales quantity per Product Brands names.



- Finally, save your report

### Save your report

Enter a name for your report \*

Select a destination workspace

LabAdmin

Save

Cancel